

Themes Paper #2

Aditi Arun Nadig

Dataset: YouTube Comments - ["AI Will Destroy Education, Teachers and Schools if we don't stop it"](#)

Overview

This paper presents the outcome of a reflexive thematic analysis I conducted on a second dataset of 303 publicly available YouTube comments posted in response to a video titled "AI Will Destroy Education, Teachers and Schools if we don't stop it." The video argues that AI in education is dangerous because it reduces learning to test score optimization, enables corporate control of schools, and removes the human teacher from the classroom. I collected comments using Outscraper.

I chose this dataset deliberately to sit in opposition to the first dataset, which came from Sal Khan's TED Talk and captured largely hopeful reactions to AI in education. The purpose was not to find a counterargument for its own sake. It was to understand what happens to people's beliefs about learning when those beliefs feel threatened. The two datasets together make visible something that neither could show alone: not just what people hope for, but what they are afraid of losing, and why that fear is so intense.

The analysis follows Braun and Clarke's (2022) reflexive thematic analysis framework, developed iteratively through familiarisation, coding, candidate theme generation, and theme revision. I also have introduced a revised research question that emerged from reading both datasets together, and a cross-dataset analysis that connects the themes from both papers through the concept of stakes.

Research Question

The original research question from Paper 1 was: how do people think AI influences education in ways that shape learning?

That question worked well for a single dataset of hopeful comments. But once Dataset 2 was collected and read, it became analytically clear that the original question could not accommodate what the data was doing. The comments here were not simply a different opinion about AI. They were angrier, more personal, and more emotionally intense than almost anything in Dataset 1. A retired professor calling the destruction of education "the real AI apocalypse." A 30-year teacher writing that without friction and effort there is "NO learning." A commenter predicting that humanity will become "helpless." These are not mild policy critiques. They are the language of people who feel something important is being taken away from them.

That observation changed the research question. The question was no longer just about what people believe. It was about what it feels like when those beliefs are threatened. Reading Dataset 2 made it clear that the beliefs from Dataset 1 - about access, emotional safety, and cognitive development - were not just hopeful ideas. They were things people felt they were owed. And the emotion in Dataset 2 was the evidence of how much it mattered to them that those things might be taken away.

The revised research question for both datasets together is: what do people's reactions to AI in education reveal about their underlying beliefs about who learning is for, and what is at stake when those beliefs are threatened?

The shift matters. The original question treats AI as the subject. The revised question treats belief as the subject, and treats AI as the trigger that made those beliefs visible and urgent enough to put into words publicly. This reframing is justified by the data itself: the intensity of the reactions in Dataset 2 could not be explained by disagreement about a technology. It could only be explained by the feeling that something deeply valued was under threat.

Positionality

My positionality for this dataset builds on what I described in Paper 1 but requires more deliberate and uncomfortable reflection, because Dataset 2 includes comments that directly challenge the position I occupy in relation to AI.

I am an immigrant student with no family support in this country. I manage everything alone - cooking, cleaning, finances, work, and school simultaneously. AI has genuinely made that survivable. It helps me structure my ideas, handle repetitive tasks, and move through the logistics of my day faster so I have capacity left for the thinking that matters. I have always been a shy student, and AI gave me a space to ask questions I was too afraid to ask in a classroom. Because I am good at asking precisely what I need, I get answers that actually satisfy the question. I also pay for the pro version of AI, which means the quality of what I get is meaningfully better than what someone on the free version receives.

I have been in tech for about four years. I lived through the era when using AI at work was penalised, and I am now in an era where recruiters actively look for people who use it well. I adapted faster than colleagues with decades more experience, partly because I entered the field without yet being attached to a particular way of doing things.

This positionality shaped my reading of Dataset 2 in ways worth naming honestly. When a 30-year programmer writes that they "wanted to cry" because AI is writing most of their code and they cannot opt out without losing their job, I understand that intellectually but I do not feel it emotionally. I never built 30 years of expertise that AI is now quietly making redundant. When commenters warn that AI will deepen inequality because the better version costs more, I sit with an uncomfortable recognition: I am on the better side of that inequality. I pay for the pro version. I know how to use it. Not everyone does. I had to keep reminding myself of that gap while coding this dataset, because my natural reading of the data was shaped by a position of relative advantage that most of these commenters do not share.

Naming this honestly is part of doing the analysis responsibly. My experience with AI is real and it has helped me. But I cannot use it as the standard when interpreting the fears and griefs of people in much more vulnerable positions.

Coding Approach

I conducted the analysis through two rounds of coding. In the first round, I applied semantic codes to the surface meaning of each comment - staying close to what commenters were directly saying. In the second round, I developed latent codes by asking what each commenter had to believe for their comment to make sense. Braun and Clarke (2022) describe this shift from semantic to latent as moving from the explicit content of the data to the assumptions and worldviews that make that content possible. From the latent codes, I developed candidate themes by grouping codes that clustered around the same central organising concept. I then revised and refined these through repeated engagement with the full dataset, checking that each theme was genuinely patterned across multiple comments rather than isolated to one or two extracts. I revised the three candidate themes produced through this process into the three finalized themes presented below.

Finalized Themes

Theme A: The Education System Was Already Broken

Central organising concept: AI did not create the crisis in education. It arrived into a system that was never truly designed for learning, and it is making that system's original failure more visible, more efficient, and harder to escape.

This theme captures something that became clear during the familiarisation phase: many commenters in Dataset 2 were not arguing against AI in isolation. They were arguing that AI is the latest arrival in a long history of things done to education in the name of improvement - standardised testing, corporate curriculum, teaching to metrics - none of which were ever really about helping students learn. The anger in these comments was not directed only at AI. It was directed at a system that had been failing long before AI arrived, and that was now finding in AI the perfect tool for continuing that failure at scale.

The comment that most precisely anchors this theme is: "It solely rewards doing what you are supposed to do. You are either punished, or completely ignored, for any learning activity that doesn't fit within its system. It does not cater to children's strengths, nor does it teach them to learn. It teaches them what to think. This has nothing to do with AI, it's just the result of a fundamentally flawed system."

The phrase "this has nothing to do with AI" is analytically significant. This commenter is not deflecting from the question. They are locating the problem precisely. AI is not the cause. It is the inheritor of a pre-existing design flaw. The belief underneath this comment is that the education system was built to produce a certain kind of person - one who follows instructions, performs on tests, and does not question the structure they are inside. That is a belief about power, not just pedagogy. It produces a specific emotional tone in the comment: not surprise that AI is being introduced in harmful ways, but weary inevitability, because this is simply what the system has always done. It finds the next tool for the same old purpose.

A second commenter adds theoretical precision to this belief: "It's Goodhart's Law - when a measure becomes a target, it ceases to be a good measure. Education has been like this for decades in the UK. AI will make it worse, for sure." Goodhart's Law is being used here not as an academic citation but as a diagnosis. The belief is that any system built around measuring performance will eventually corrupt the thing it is measuring. Test scores stopped representing learning the moment they became the goal. AI, introduced specifically to improve test scores, will accelerate that corruption - not because AI is uniquely dangerous, but because it is being handed a broken brief.

A third comment captures the emotional weight of this theme: "Within the system, hardly anyone cares about actually teaching or learning something. It's all just an oiled machine to produce non-thinking obedient wage slaves." The phrase "oiled machine" is doing important work. A machine does not see individuals. It does not adapt. It does not notice when something is wrong with a student's life or circumstances. This commenter is saying that school has always functioned this way - with indifference to the individual baked into its design. Adding AI does not change that design. It makes the machine run more efficiently and at greater scale. Whatever small moments of genuine human connection existed inside classrooms - a teacher who noticed a struggling student, a moment of genuine encouragement - those are exactly what gets removed when the human is replaced by a system optimising for measurable outputs.

My positionality shaped how I read this theme. As an immigrant student who entered a school system not designed for me - one that assumed cultural familiarity, family support, and a particular prior knowledge I did not have - I recognised the experience being described here. The system not seeing you is not an abstraction. That personal recognition made me more attuned to the nuances of this theme, but it also meant I had to be

careful not to project my own specific experience onto comments that were sometimes describing a different kind of institutional failure.

Theme B: AI Takes Away the Ability to Think

Central organising concept: real learning happens through effort, struggle, and the discomfort of not knowing. AI removes that friction. A generation that grows up without that friction may lose the capacity to think independently entirely.

This theme is the most direct continuation of a conversation that began in Dataset 1. Theme 3 from Paper 1 captured a cautious, uncertain worry - commenters sensing that AI might remove the productive struggle that makes learning stick, but uncertain about how serious the consequences would be. Dataset 2 commenters have resolved that uncertainty. They are not speculating about a future risk. They are reporting something they believe they have already observed, and they are alarmed by it.

The comment that anchored this theme analytically came from a 30-year teacher: "AI is a frictionless tool that smooths the path for students to complete assignments and assessments. Without friction, without effort and without intention there is NO learning. Students are turning in products rather than work."

The word "frictionless" is doing precise analytic work here. Friction is what slows things down, forces engagement, and demands that you actually think. A frictionless path to a completed assignment is not learning - it is production. The teacher is drawing a distinction between work as a process and work as a product. AI delivers the product without the process. And if the process is where learning happens - if struggle is the mechanism, not just the unfortunate side effect - then the product without the process is an empty shell. The capital letters on "NO learning" are not stylistic. They signal emphasis and certainty. This teacher is not offering a hypothesis. They are reporting something they believe they have watched happen across three decades of working with students, and they are certain about what they are seeing.

A second comment extends this concern to a societal level: "We are training an entire generation of kids that critical independent thought can be outsourced. There is no free lunch." The phrase "no free lunch" is economically precise. Every apparent shortcut has a hidden cost. The cost of outsourcing thinking is the gradual atrophy of the capacity to think. This commenter is not worried about individual students failing their assignments. They are worried about what an entire generation that grew up outsourcing cognition will look like when they face decisions that have no AI prompt waiting to help them.

The comment I found most analytically rich for this theme came from a programmer: "In my nearly 30-year career, I have never felt so detached from the application code my team has been writing. If I don't embrace it I stand to lose my job." What makes this comment significant is the last sentence. This person is not using AI by choice. They are using it because refusing means losing their livelihood. That distinction - between engaging with AI freely and engaging under compulsion - reveals something important about the nature of the threat. It is not just that AI changes how people think. It is that it is removing the choice of whether to engage with it at all. The grief in this comment is not about a tool. It is about the loss of agency over one's own cognitive practice and professional identity.

A third comment brings the stakes to their most extreme articulation: "We are about to offload all our cognitive ability to a textbox and a prompt. Our creativity will be gone. Decision-making gone. Ability to evaluate gone. And then, inevitably, the systems will fail and most of the world will starve because we've become helpless." This reads as extreme, but the analytic structure underneath it is serious. The commenter is making a systems-level argument: if enough people stop developing their own cognitive capacity, collective intelligence degrades. If the systems on which everyone depends then fail, there will not be enough people who retained

independent thinking to hold things together. The word "helpless" is the key. It describes people who want to act but have lost the capacity to do so without external assistance. That is not a comment about a bad app. It is a comment about what happens to a civilisation that stops requiring its members to think.

My positionality is most complicated in relation to this theme. I use AI constantly and I do not experience a loss of thinking. I generate original ideas and use AI to bring structure to them. The thinking is mine. But I entered tech four years ago and developed foundational skills before AI was capable of doing everything for me. I do not know whether my thinking would be different - more rigorous, more independent - if AI had been unavailable during those years. That uncertainty is worth naming honestly, because it shaped how much emotional weight I was able to give to the fear these comments express.

Theme C: AI in Education Serves Corporations, Not Students

Central organising concept: AI is not being introduced into schools by educators, communities, or students. It is being introduced by corporations with financial interests. The result will not be equal access to better learning. It will be a tiered market in which the quality of education a child receives depends on what their family can pay - and whoever controls that market controls what the next generation knows, thinks, and believes.

This theme sits in direct opposition to Theme 1 from Paper 1, which captured the hopeful belief that AI could finally democratise education. Dataset 2 commenters look at the same technology and see the opposite happening.

The comment that most fully captures this theme is: "Those who can afford the 'super educator 3000' model might have a chance of getting respectable education, further widening the inequality gap. Once you control education you control society."

The phrase "super educator 3000" reads like a product name - something packaged and marketed with large promises attached. The framing is analytically revealing regardless of whether it was intended as mockery. The commenter is describing AI education as a product, not a public service. And products have tiers. The expensive version works better. The cheap version is optimised for something else - in this case, probably test scores and user retention rather than genuine learning. The latent belief in the first sentence is that the same inequality that has always existed in education will simply be rebuilt inside AI, with a premium tier for those who can pay and a stripped-down version for everyone else. That is not a speculation about the future. For this commenter, it is a prediction based on a pattern they have already observed in how markets work.

The second sentence carries the deepest analytical weight: "Once you control education you control society." This is not a complaint about a bad product. It is a belief about power. Whoever controls what children learn controls what they grow up to think, value, and accept. If that control sits with a corporation whose primary accountability is to shareholders rather than to students or communities, the consequences extend far beyond whether the AI tutor delivers good pedagogy. It means the cognitive habits, knowledge base, and belief systems of the next generation are being shaped by a profit motive. The emotion underneath this comment is not irritation at a bad product. It is fear about who is being handed control of the future.

A second comment proves analytically difficult to set aside: "The folks who got rich with social media sent their kids to no-smartphones-allowed schools. The folks who get rich with AI will send their kids to no-AI schools. You don't give crack to your kids, you sell it to idiots." The bluntness is part of its analytic value. If the people building and profiting from AI education tools are simultaneously protecting their own children from those tools, the product being sold to everyone else is not genuinely good. It is being sold to people who either cannot afford to say no, or do not have enough information to know they should. The underlying belief is about information asymmetry - the sellers know something the buyers do not. That asymmetry maps directly onto

existing economic inequality: the families with resources and information can protect their children from AI, while the families without resources accept it because they have no alternative.

A third comment addresses the question of institutional power: "If you are dependent on a single AI company to provide your teachers then you are at the mercy of that company which will drive a much harder bargain than a collection of human workers ever could." This commenter is describing a transfer of institutional leverage. A corporation that becomes the sole provider of education holds a kind of structural power that no labour union or government body could easily counteract. The belief underneath is that introducing AI into schools does not just change how students learn. It changes who holds power over the conditions of learning - shifting that power away from communities, teachers, and public institutions and into the hands of a private entity with no public accountability.

My positionality is most uncomfortable here. I am an immigrant student without a family support network, which places me in the group these commenters are warning about. But I pay for the pro version of AI and know how to use it effectively, which places me in the advantaged group. The tiering these commenters describe already exists. I happen to be on the better side of it. That does not make the warning wrong. It means I had to be especially deliberate about not generalising from my own experience when interpreting the experience of people with fewer resources and less domain knowledge. The reflexive discipline required here was not just intellectual. It was genuinely uncomfortable to recognise how much of my ease with AI depends on access and skills that are themselves unevenly distributed.

Cross-Dataset Analysis: What Is at Stake

The three themes from Dataset 2 do not just contrast with the three themes from Dataset 1. They respond to them directly. Each pairing reveals a stake - something people believed education was supposed to give them, which is now at risk. The stakes only become visible when both datasets are present. You cannot see what someone is protecting until something threatens to take it away. Dataset 2 is that threat. And the intensity of the reactions - the anger, the grief, the alarm - tells you how much was at stake in Dataset 1 all along.

Stake 1: The hope that education could finally be fair

Theme 1 from Paper 1 captured the belief that education has always been unequal and AI might finally change that. Commenters described AI as a route in for people locked out by cost, geography, disability, and privilege. For many, this was not abstract. It was personal. "Learning will now be accessible for the forgotten, misplaced, and abandoned." The words "forgotten, misplaced, and abandoned" are not neutral. They describe people who were not just unlucky but actively failed by a system that was supposed to serve them. The hope in this comment is enormous - and it is the hope of someone who has likely experienced that failure firsthand.

Theme C directly dismantles it. The same inequality will simply be rebuilt inside AI. Premium tools for those who can pay, a cheaper and less effective version for everyone else. The product being marketed as democratisation is actually the foundation of a new hierarchy. The stake between these two themes is the hope that education could finally be fair. Dataset 1 holds that hope. Dataset 2 warns it will not be delivered.

Stake 2: The belief that learning can feel human and caring

Theme 2 from Paper 1 captured the belief that good learning requires emotional safety - a space where you can ask questions without fear, where the system adapts to you rather than demanding you adapt to it. "It doesn't get angry and lets the student ask questions he is afraid to ask in class with 20 other students making fun of him." The specificity of this image - the 20 other students, the mockery, the anger of the teacher -

suggests this commenter is describing something they have lived. The belief is that emotional safety is a precondition for learning. You cannot think when you are ashamed or afraid.

Theme A responds by arguing the system was never human or caring to begin with. It was built like a factory, and AI will not make it warmer. It will make the factory more efficient. The stake here is the belief that learning can ever feel genuinely human. Dataset 1 commenters found that in AI. Dataset 2 commenters warn that what AI is being introduced into will ensure that experience remains an exception rather than a standard.

Stake 3: The belief that thinking is a skill built through struggle

Theme 3 from Paper 1 captured a cautious worry that AI might remove the productive struggle that makes learning stick. Commenters were uncertain - they could see AI's value while sensing something important might be lost. "Isn't it good to sit with a problem and be frustrated?" The question format matters here. This commenter is not certain. They are genuinely asking, which means they are still working out what they believe. That uncertainty is the emotional register of Dataset 1 on this question.

Theme B removes all uncertainty. Dataset 2 commenters are not asking. They are reporting. "Without friction, without effort and without intention there is NO learning." The shift from uncertain question to certain declaration is the emotional difference between the two datasets on this stake, and it is analytically significant. The stake is the belief that thinking is something you earn through effort and can lose through disuse. Dataset 1 worries about this possibility. Dataset 2 grieves it as something already happening.

Conclusion

This project began with a question about AI and education. It ended with an answer about what people believe learning was always supposed to give them, and what they feel when that starts to slip away.

People are not arguing about AI. They are defending three things they believed education owed them: equal access, human connection, and the ability to think. Those beliefs were present in Dataset 1 as hope. They are present in Dataset 2 as fear and grief. The difference in emotional register between the two datasets is not a methodological problem. It is the most important finding of the analysis. It shows that what people feel about AI in education is proportional to how much they believed education was supposed to deliver in the first place - and how far short of that it has always fallen.

The revised research question asks what people's reactions reveal about their underlying beliefs and what is at stake when those beliefs are threatened. The answer is: the stakes are everything people believed learning was supposed to be. And the intensity of the response in Dataset 2 is the clearest evidence of how much those beliefs have always mattered, even when, especially when, the system failed to honour them.

Data

[Dataset #1](#)

[Dataset #2](#)

[Mural Board](#)